


Informer: Beyond Efficient Transformer for Long Sequence Time-Series Forecasting

Zhou, H., Zhang, S., Peng, J., Zhang, S., Li, J., Xiong, H., & Zhang, W. (2021, May).
In *Proceedings of the AAAI conference on artificial intelligence*
(Vol. 35, No. 12, pp. 11106-11115).

- 01 Background
- 02 Methodology
- 03 Experiment
- 04 Conclusion
- 05 Final proposal

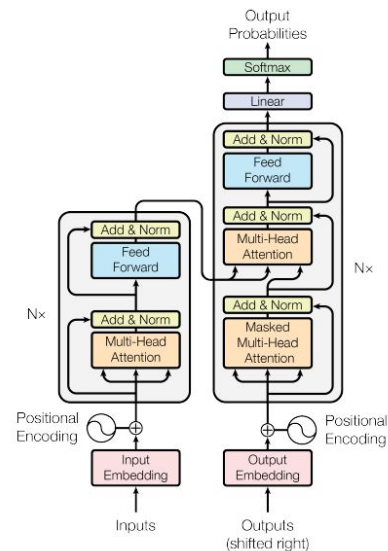
Background

- Long sequence time-series forecasting (LSTF) is crucial across many domains
- How to enhance the prediction capacity of LSTF?
 - Long-range alignment ability
 - Efficient operations on long sequence inputs and outputs



Problems of Transformer


- The quadratic computation of self-attention $O(L^2)$
→ ProbSparse self-attention $O(L\log L)$
- The memory bottleneck in stacking layers for long inputs
→ Self-attention distilling operation
- The speed plunge in predicting long outputs
→ Generative-style decoder



Methodology


How does self-attention work?

- Self-attention: Consider the context to get relevance.
 - [NIPS 2017] Attention is all you need!
- Give a sequence of data, return a sequence of corresponding attention score.




How does self-attention work?

- The canonical self-attention is defined based on the tuple inputs.
 - query, key and value (That is, we have to learn the weight W^q , W^k , W^v)



How does self-attention work?

- Performs the scaled dot-product as $A(Q, K, V) = \text{Softmax}\left(\frac{QK^T}{\sqrt{d}}\right)V$ **Problem**
 - where $Q \in \mathbb{R}^{L_Q \times d}$, $K \in \mathbb{R}^{L_K \times d}$, $V \in \mathbb{R}^{L_V \times d}$ and d is the input dimension

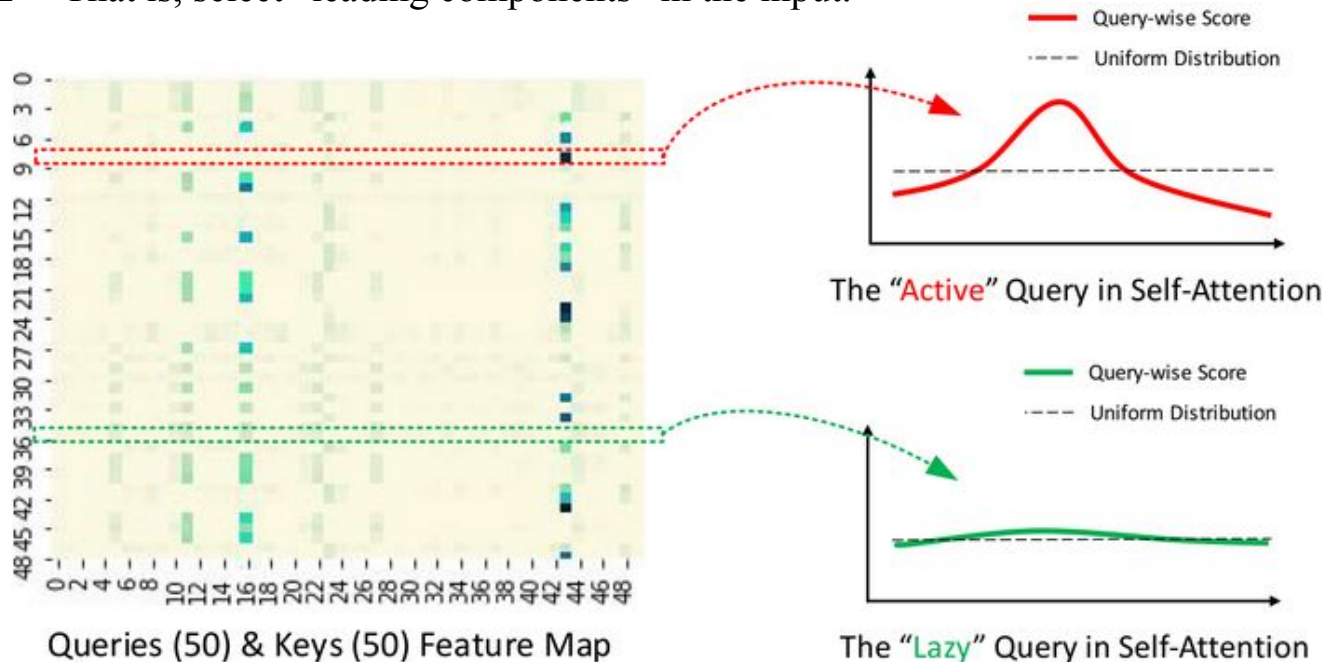


time & space complexity: $O(L^2)$

$$\begin{aligned}
 \begin{matrix} b^1 & b^2 & b^3 & b^4 \\ \hline 0 \end{matrix} &= \begin{matrix} v^1 & v^2 & v^3 & v^4 \\ \hline V \end{matrix} \\
 &= \begin{matrix} \alpha'_{1,1} & \alpha'_{2,1} & \alpha'_{3,1} & \alpha'_{4,1} \\ \alpha'_{1,2} & \alpha'_{2,2} & \alpha'_{3,2} & \alpha'_{4,2} \\ \alpha'_{1,3} & \alpha'_{2,3} & \alpha'_{3,3} & \alpha'_{4,3} \\ \alpha'_{1,4} & \alpha'_{2,4} & \alpha'_{3,4} & \alpha'_{4,4} \end{matrix} \xleftarrow{\text{softmax}} \begin{matrix} \alpha_{1,1} & \alpha_{2,1} & \alpha_{3,1} & \alpha_{4,1} \\ \alpha_{1,2} & \alpha_{2,2} & \alpha_{3,2} & \alpha_{4,2} \\ \alpha_{1,3} & \alpha_{2,3} & \alpha_{3,3} & \alpha_{4,3} \\ \alpha_{1,4} & \alpha_{2,4} & \alpha_{3,4} & \alpha_{4,4} \end{matrix} = \begin{matrix} k^1 \\ k^2 \\ k^3 \\ k^4 \\ \hline K^T \end{matrix} \begin{matrix} q^1 & q^2 & q^3 & q^4 \\ \hline Q \end{matrix} \\
 & \quad \quad \quad A' \quad \quad \quad A
 \end{aligned}$$

Query Sparsity Measurement

- Self-attention probability has potential **sparsity** and also **long tail distribution**.
 - We want to find out “Active” Query.
 - That is, select “leading components” in the input.



Query Sparsity Measurement

- How to find out dominant dot product pairs
 - Use KL-divergence.

$$\mathcal{A}(\mathbf{q}_i, \mathbf{K}, \mathbf{V}) = \sum_j \frac{k(\mathbf{q}_i, \mathbf{k}_j)}{\sum_l k(\mathbf{q}_i, \mathbf{k}_l)} \mathbf{v}_j = \mathbb{E}_{p(\mathbf{k}_j|\mathbf{q}_i)}[\mathbf{v}_j]$$

the i-th query's attention on all the keys


$$q(\mathbf{k}_j|\mathbf{q}_i) = 1/L_K \leftarrow \text{uniform distribution}$$

$$KL(q||p) = \ln \sum_{l=1}^{L_K} e^{\mathbf{q}_i \mathbf{k}_l^\top / \sqrt{d}} - \frac{1}{L_K} \sum_{j=1}^{L_K} \mathbf{q}_i \mathbf{k}_j^\top / \sqrt{d} - \ln L_K$$

Get Top-u queries to calculate self-attention

How to reduce time & space complexity?

- Calculate all dot-product pairs to find “active query”. → $O(L^2)$ **NO!!**
- Since long tail distribution, we only need to randomly sample $c \cdot \ln L$ keys
 - Active query is high-correlated with other keys.
 - Authors investigate vanilla Transformer on ETT dataset



sample $c \cdot \ln L$ keys

Algorithm 1 ProbSparse self-attention

Require: Tensor $\mathbf{Q} \in \mathbb{R}^{m \times d}$, $\mathbf{K} \in \mathbb{R}^{n \times d}$, $\mathbf{V} \in \mathbb{R}^{n \times d}$

1: **print** set hyperparameter c , $u = c \ln m$ and $U = m \ln n$

2: randomly select U dot-product pairs from \mathbf{K} as $\bar{\mathbf{K}}$

$O(L \ln L)$ 3: set the sample score $\bar{\mathbf{S}} = \mathbf{Q}\bar{\mathbf{K}}^\top$

4: compute the measurement $M = \max(\bar{\mathbf{S}}) - \text{mean}(\bar{\mathbf{S}})$ by row

5: set Top- u queries under M as $\bar{\mathbf{Q}}$


$O(L \ln L)$ 6: set $\mathbf{S}_1 = \text{softmax}(\bar{\mathbf{Q}}\mathbf{K}^\top / \sqrt{d}) \cdot \mathbf{V}$

7: set $\mathbf{S}_0 = \text{mean}(\mathbf{V})$

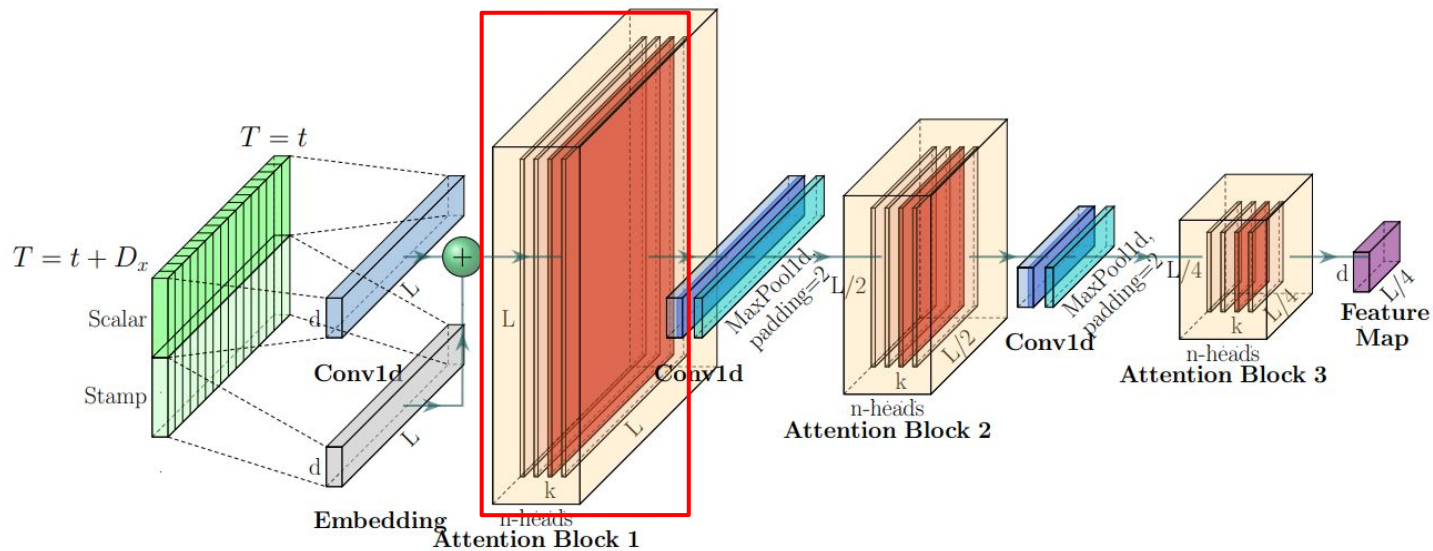
8: set $\mathbf{S} = \{\mathbf{S}_1, \mathbf{S}_0\}$ by their original rows accordingly

Ensure: self-attention feature map \mathbf{S} .


- Input



- ProbSparse self-attention



- Distilling




- Stack more **attention blocks** can get more detail features.
- Problem: the memory usage is $O(J \cdot L^2)$ if stack J layers
- Solution: Distilling operation

Self Attention Distilling

- Distilling operation to privilege the superior ones with dominating features

$$\mathbf{X}_{j+1}^t = \text{MaxPool} \left(\text{ELU} \left(\text{Conv1d}([\mathbf{X}_j^t]_{AB}) \right) \right)$$




Self Attention Distilling

Active Query: $\mathbf{S}_1 = \text{softmax}(\bar{\mathbf{Q}}\mathbf{K}^\top / \sqrt{d}) \cdot \mathbf{V}$

Lazy Query: $\mathbf{S}_0 = \text{mean}(\mathbf{V})$


$\mathbf{S} = \{\mathbf{S}_1, \mathbf{S}_0\}$ by their original rows accordingly

- ❖ Lazy Queries are not important.
- Use maxpooling to make active queries dominate features map.
 - However, it lost some information.




$$\mathbf{X}_{j+1}^t = \text{MaxPool} \left(\text{ELU} \left(\text{Conv1d}([\mathbf{X}_j^t]_{AB}) \right) \right)$$

Layer 2



reduce

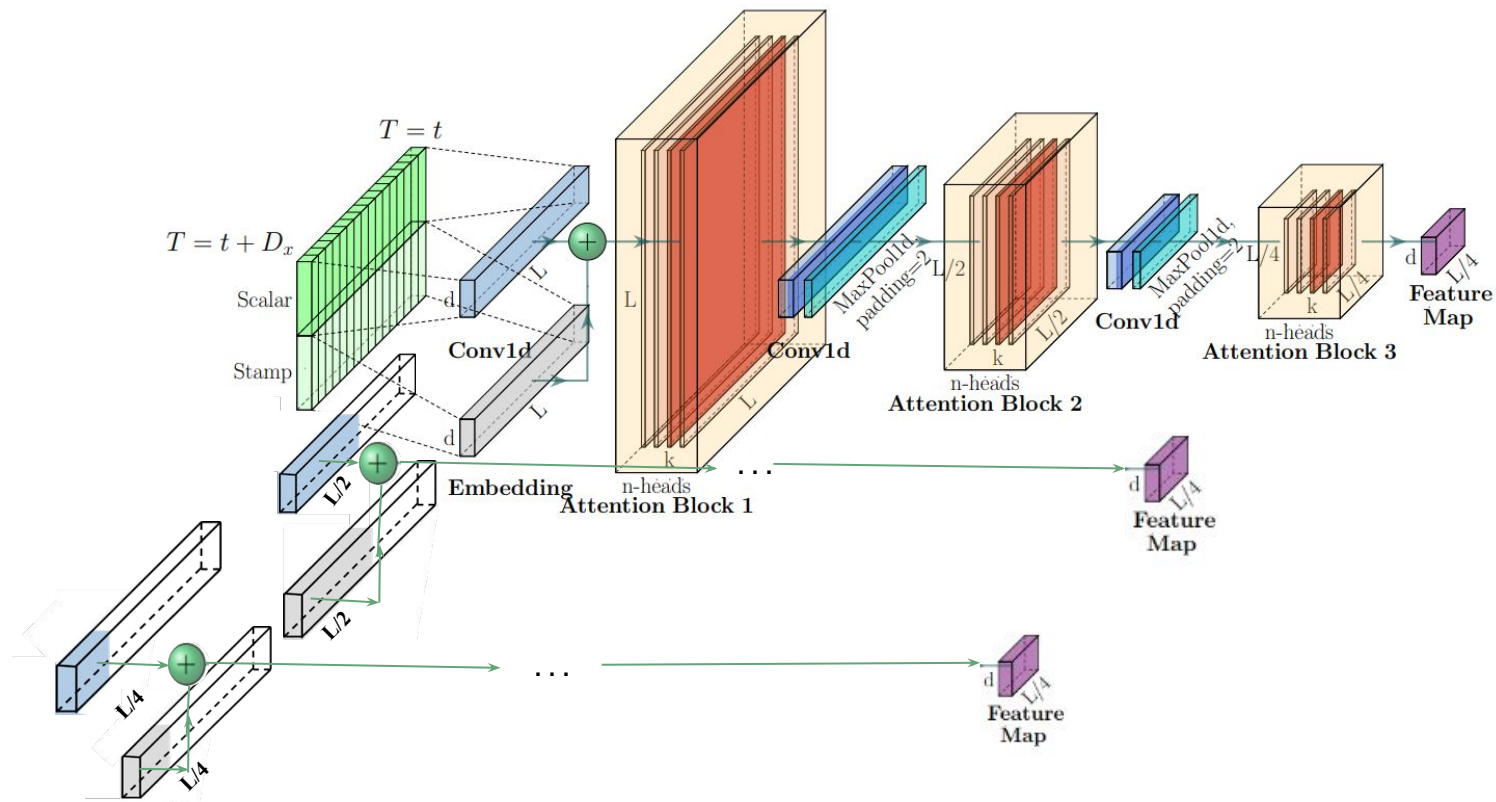
Layer 2



CSDN@郑烯怪快去学习

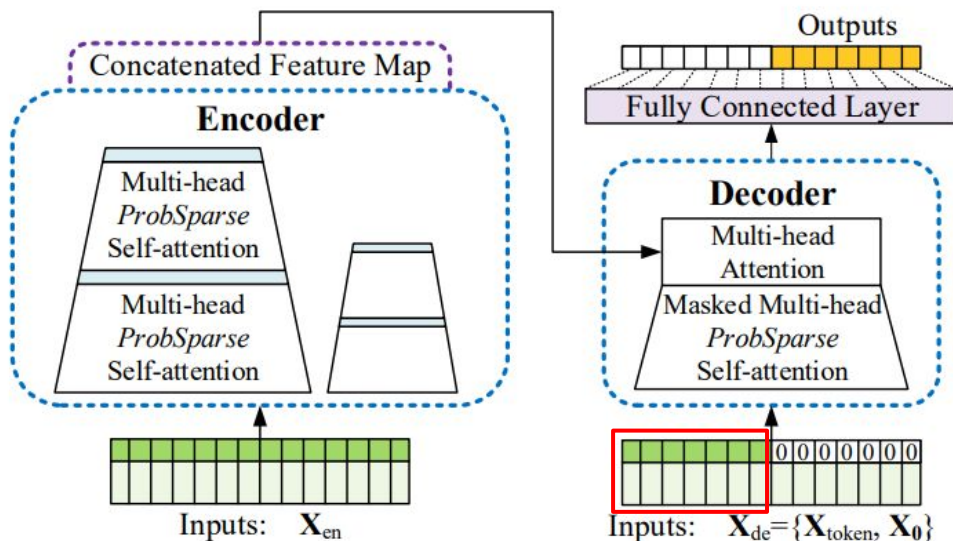
Self Attention Distilling

- Enhance the robustness of the distilling operation



- One Forward Procedure → Use shorter input sequence(groundtruth) instead of specific flag as “start token” in decoder.

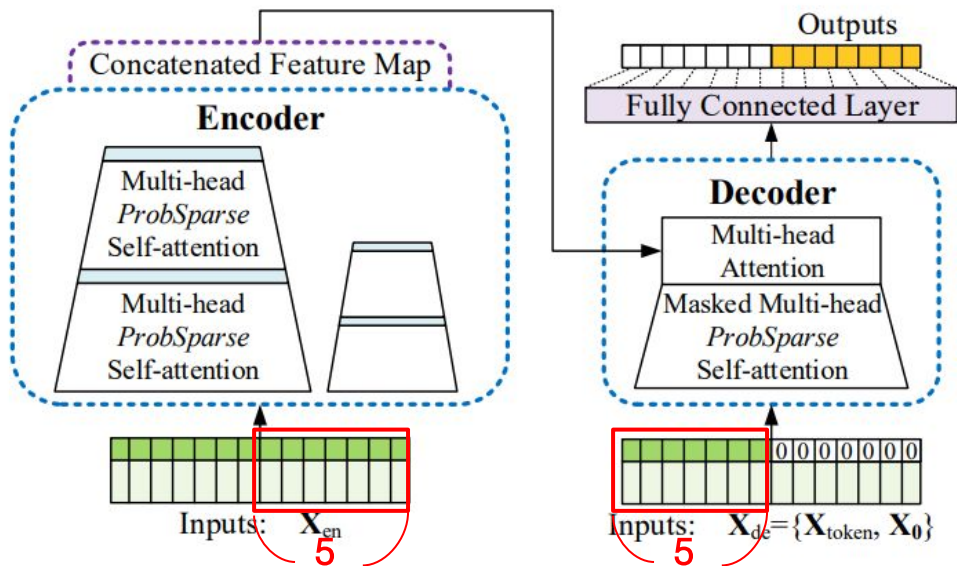
$$\mathbf{X}_{\text{de}}^t = \text{Concat}(\mathbf{X}_{\text{token}}^t, \mathbf{X}_0^t) \in \mathbb{R}^{(L_{\text{token}} + L_y) \times d_{\text{model}}}$$



- One Forward Procedure

$$\mathbf{X}_{\text{de}}^t = \text{Concat}(\mathbf{X}_{\text{token}}^t, \mathbf{X}_0^t) \in \mathbb{R}^{(L_{\text{token}} + L_y) \times d_{\text{model}}}$$

Predict 7-day temperature, x_{token}^t = the 5 days before the forecasted time point.



Experiment

Experiment

1. Electricity Transformer Temperature(ETT)
2. Electricity Consuming Load (ECL)
3. Weather

Station	Date	Latitude	Longitude	SeaLevelPressure	WindDirection	WindSpeed	WetBulbTemperature
2907099999	1903-01-01T08:00:00	64.3333333	23.45	29.8	90	9	3
2907099999	1903-01-01T15:00:00	64.3333333	23.45	29.78	90	14	4
2907099999	1903-01-01T22:00:00	64.3333333	23.45	29.65	50	26	10
2907099999	1903-01-02T08:00:00	64.3333333	23.45	29.61	360	45	13
2907099999	1903-01-02T15:00:00	64.3333333	23.45	29.65	340	36	11

...

- All dataset will undergo **univariate** and **multivariate** time-series forecasting.

Univariate


Methods		Informer		Informer [†]		LogTrans		Reformer		LSTMa		DeepAR		ARIMA		Prophet	
Metric		MSE	MAE	MSE	MAE	MSE	MAE	MSE	MAE	MSE	MAE	MSE	MAE	MSE	MAE	MSE	MAE
ETTh ₁	24	0.098	0.247	0.092	0.246	0.103	0.259	0.222	0.389	0.114	0.272	0.107	0.280	0.108	0.284	0.115	0.275
	48	0.158	0.319	0.161	0.322	0.167	0.328	0.284	0.445	0.193	0.358	0.162	0.327	0.175	0.424	0.168	0.330
	168	0.183	0.346	0.187	0.355	0.207	0.375	1.522	1.191	0.236	0.392	0.239	0.422	0.396	0.504	1.224	0.763
	336	0.222	0.387	0.215	0.369	0.230	0.398	1.860	1.124	0.590	0.698	0.445	0.552	0.468	0.593	1.549	1.820
	720	0.269	0.435	0.257	0.421	0.273	0.463	2.112	1.436	0.683	0.768	0.658	0.707	0.659	0.766	2.735	3.253
ETTh ₂	24	0.093	0.240	0.099	0.241	0.102	0.255	0.263	0.437	0.155	0.307	0.098	0.263	3.554	0.445	0.199	0.381
	48	0.155	0.314	0.159	0.317	0.169	0.348	0.458	0.545	0.190	0.348	0.163	0.341	3.190	0.474	0.304	0.462
	168	0.232	0.389	0.235	0.390	0.246	0.422	1.029	0.879	0.385	0.514	0.255	0.414	2.800	0.595	2.145	1.068
	336	0.263	0.417	0.258	0.423	0.267	0.437	1.668	1.228	0.558	0.606	0.604	0.607	2.753	0.738	2.096	2.543
	720	0.277	0.441	0.285	0.442	0.303	0.493	2.030	1.721	0.640	0.681	0.629	0.580	2.878	1.044	3.355	4.664
ETTm ₁	24	0.030	0.137	0.034	0.160	0.065	0.202	0.095	0.228	0.121	0.233	0.091	0.243	0.090	0.206	0.120	0.290
	48	0.069	0.203	0.066	0.194	0.078	0.220	0.249	0.390	0.305	0.411	0.219	0.362	0.179	0.306	0.133	0.305
	96	0.194	0.372	0.187	0.384	0.199	0.386	0.920	0.767	0.287	0.420	0.364	0.496	0.272	0.399	0.194	0.396
	288	0.401	0.554	0.409	0.548	0.411	0.572	1.108	1.245	0.524	0.584	0.948	0.795	0.462	0.558	0.452	0.574
	672	0.512	0.644	0.519	0.665	0.598	0.702	1.793	1.528	1.064	0.873	2.437	1.352	0.639	0.697	2.747	1.174
Weather	24	0.117	0.251	0.119	0.256	0.136	0.279	0.231	0.401	0.131	0.254	0.128	0.274	0.219	0.355	0.302	0.433
	48	0.178	0.318	0.185	0.316	0.206	0.356	0.328	0.423	0.190	0.334	0.203	0.353	0.273	0.409	0.445	0.536
	168	0.266	0.398	0.269	0.404	0.309	0.439	0.654	0.634	0.341	0.448	0.293	0.451	0.503	0.599	2.441	1.142
	336	0.297	0.416	0.310	0.422	0.359	0.484	1.792	1.093	0.456	0.554	0.585	0.644	0.728	0.730	1.987	2.468
	720	0.359	0.466	0.361	0.471	0.388	0.499	2.087	1.534	0.866	0.809	0.499	0.596	1.062	0.943	3.859	1.144
ECL	48	0.239	0.359	0.238	0.368	0.280	0.429	0.971	0.884	0.493	0.539	0.204	0.357	0.879	0.764	0.524	0.595
	168	0.447	0.503	0.442	0.514	0.454	0.529	1.671	1.587	0.723	0.655	0.315	0.436	1.032	0.833	2.725	1.273
	336	0.489	0.528	0.501	0.552	0.514	0.563	3.528	2.196	1.212	0.898	0.414	0.519	1.136	0.876	2.246	3.077
	720	0.540	0.571	0.543	0.578	0.558	0.609	4.891	4.047	1.511	0.966	0.563	0.595	1.251	0.933	4.243	1.415
	960	0.582	0.608	0.594	0.638	0.624	0.645	7.019	5.105	1.545	1.006	0.657	0.683	1.370	0.982	6.901	4.264
Count		32		12		0		0		0		6		0		0	

Multivariate

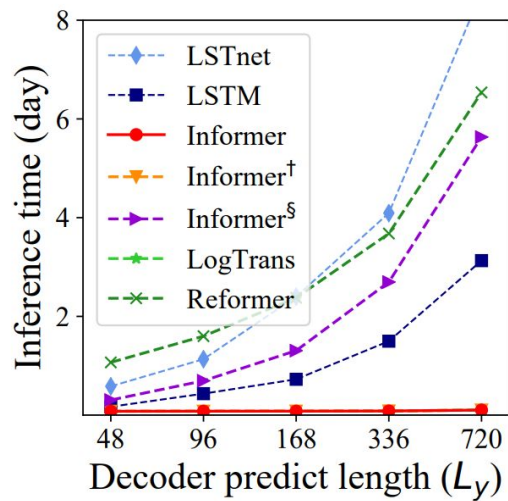
Methods		Informer		Informer [†]		LogTrans		Reformer		LSTMa		LSTNet	
Metric		MSE	MAE	MSE	MAE	MSE	MAE	MSE	MAE	MSE	MAE	MSE	MAE
ETTh ₁	24	0.577	0.549	0.620	0.577	0.686	0.604	0.991	0.754	0.650	0.624	1.293	0.901
	48	0.685	0.625	0.692	0.671	0.766	0.757	1.313	0.906	0.702	0.675	1.456	0.960
	168	0.931	0.752	0.947	0.797	1.002	0.846	1.824	1.138	1.212	0.867	1.997	1.214
	336	1.128	0.873	1.094	0.813	1.362	0.952	2.117	1.280	1.424	0.994	2.655	1.369
	720	1.215	0.896	1.241	0.917	1.397	1.291	2.415	1.520	1.960	1.322	2.143	1.380
ETTh ₂	24	0.720	0.665	0.753	0.727	0.828	0.750	1.531	1.613	1.143	0.813	2.742	1.457
	48	1.457	1.001	1.461	1.077	1.806	1.034	1.871	1.735	1.671	1.221	3.567	1.687
	168	3.489	1.515	3.485	1.612	4.070	1.681	4.660	1.846	4.117	1.674	3.242	2.513
	336	2.723	1.340	2.626	1.285	3.875	1.763	4.028	1.688	3.434	1.549	2.544	2.591
	720	3.467	1.473	3.548	1.495	3.913	1.552	5.381	2.015	3.963	1.788	4.625	3.709
ETTm ₁	24	0.323	0.369	0.306	0.371	0.419	0.412	0.724	0.607	0.621	0.629	1.968	1.170
	48	0.494	0.503	0.465	0.470	0.507	0.583	1.098	0.777	1.392	0.939	1.999	1.215
	96	0.678	0.614	0.681	0.612	0.768	0.792	1.433	0.945	1.339	0.913	2.762	1.542
	288	1.056	0.786	1.162	0.879	1.462	1.320	1.820	1.094	1.740	1.124	1.257	2.076
	672	1.192	0.926	1.231	1.103	1.669	1.461	2.187	1.232	2.736	1.555	1.917	2.941
Weather	24	0.335	0.381	0.349	0.397	0.435	0.477	0.655	0.583	0.546	0.570	0.615	0.545
	48	0.395	0.459	0.386	0.433	0.426	0.495	0.729	0.666	0.829	0.677	0.660	0.589
	168	0.608	0.567	0.613	0.582	0.727	0.671	1.318	0.855	1.038	0.835	0.748	0.647
	336	0.702	0.620	0.707	0.634	0.754	0.670	1.930	1.167	1.657	1.059	0.782	0.683
	720	0.831	0.731	0.834	0.741	0.885	0.773	2.726	1.575	1.536	1.109	0.851	0.757
ECL	48	0.344	0.393	0.334	0.399	0.355	0.418	1.404	0.999	0.486	0.572	0.369	0.445
	168	0.368	0.424	0.353	0.420	0.368	0.432	1.515	1.069	0.574	0.602	0.394	0.476
	336	0.381	0.431	0.381	0.439	0.373	0.439	1.601	1.104	0.886	0.795	0.419	0.477
	720	0.406	0.443	0.391	0.438	0.409	0.454	2.009	1.170	1.676	1.095	0.556	0.565
	960	0.460	0.548	0.492	0.550	0.477	0.589	2.141	1.387	1.591	1.128	0.605	0.599
Count		33		14		1		0		0		2	

- Time consumption

Training time:



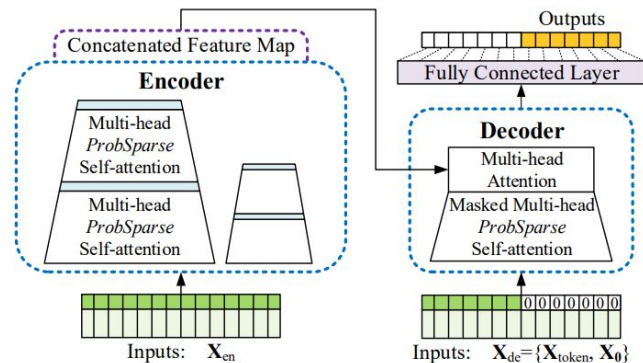
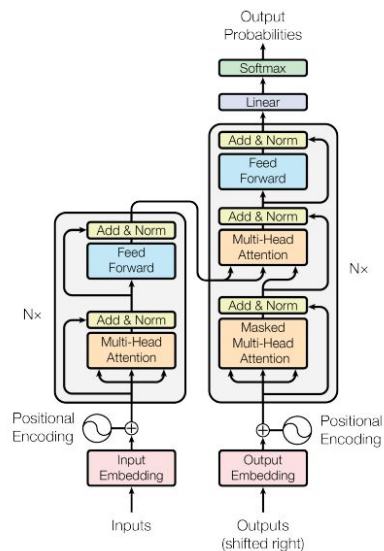
Inference time:



Conclusion

Conclusion

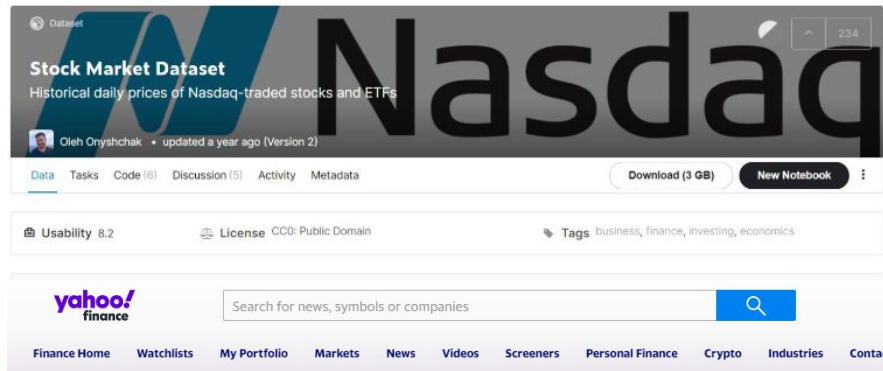
- Reduce time complexity \rightarrow ProbSparse self-attention
- Reduce memory usage \rightarrow ProbSparse self-attention
& Self-attention distilling operation
- Reduce inference time \rightarrow Generative style decoder




Final proposal

- Dataset: Stock
- We want to predict companies' closing returns given stock features.
 - Date
 - Open
 - High
 - Low
 - Closing
 - Volume
- Univariate

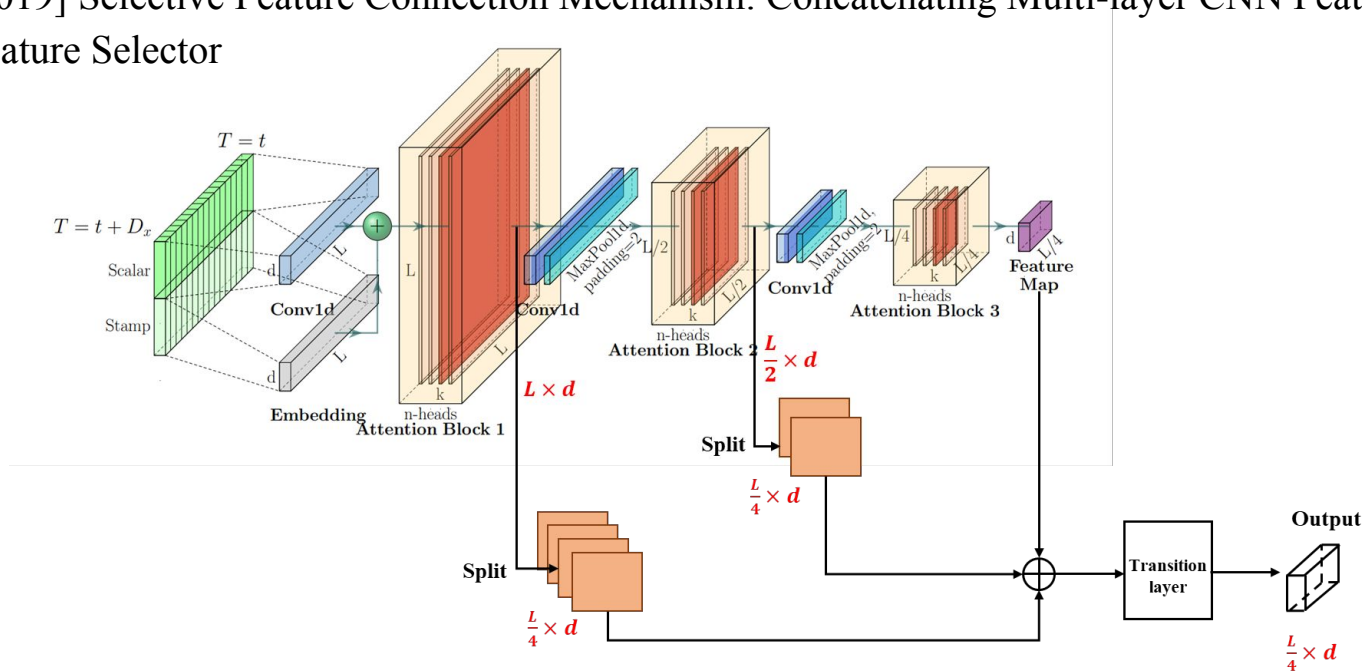
Date	Open	High	Low	Close	Adj Close	Volume
2020/2/6	30.99	31.56	29.56	30	30	2552630
2020/2/7	29.75	31.75	29.71	30.92	30.92	357500
2020/2/10	31.8	32	31	31.89	31.89	229510
2020/2/11	31.94	33.23	31.93	32.87	32.87	286300



- Problem: How to determine replicas of the main stack with halving inputs?
 - the first half, the second half or random selection?

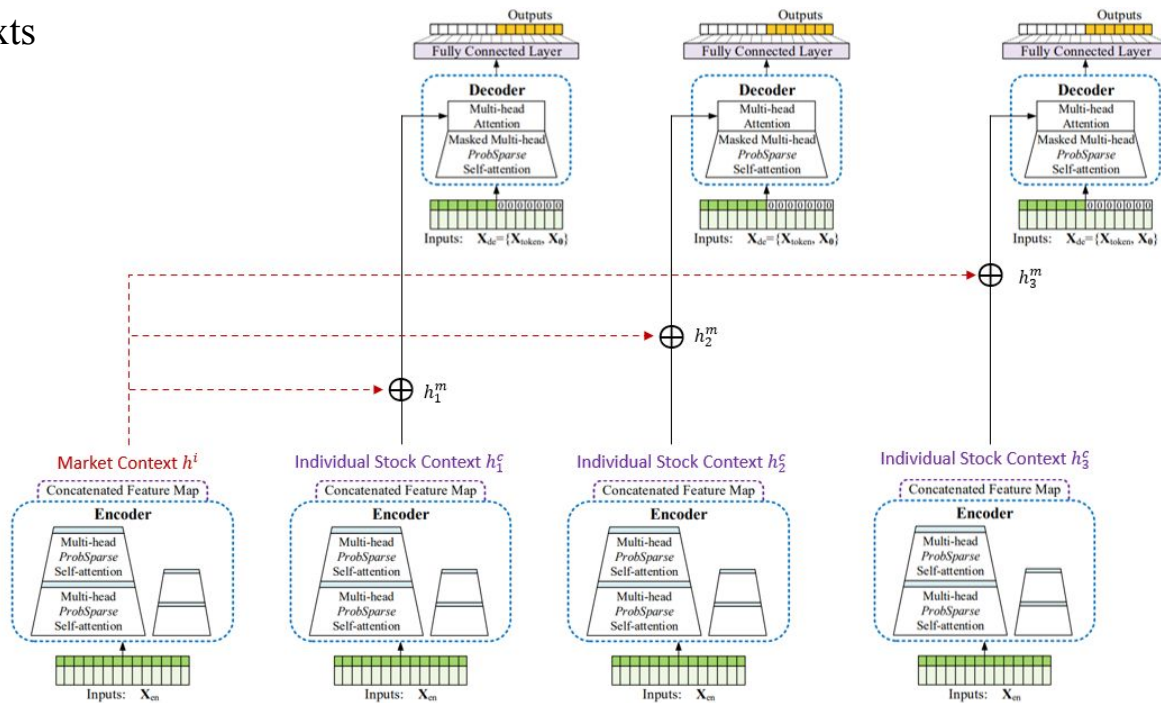


- Feature combination of the low-layer and high-level features
 - High-layer features contain more semantic information.
 - Low-layer features contain more detail information.
- ❖ [2019] Selective Feature Connection Mechanism: Concatenating Multi-layer CNN Features with a Feature Selector

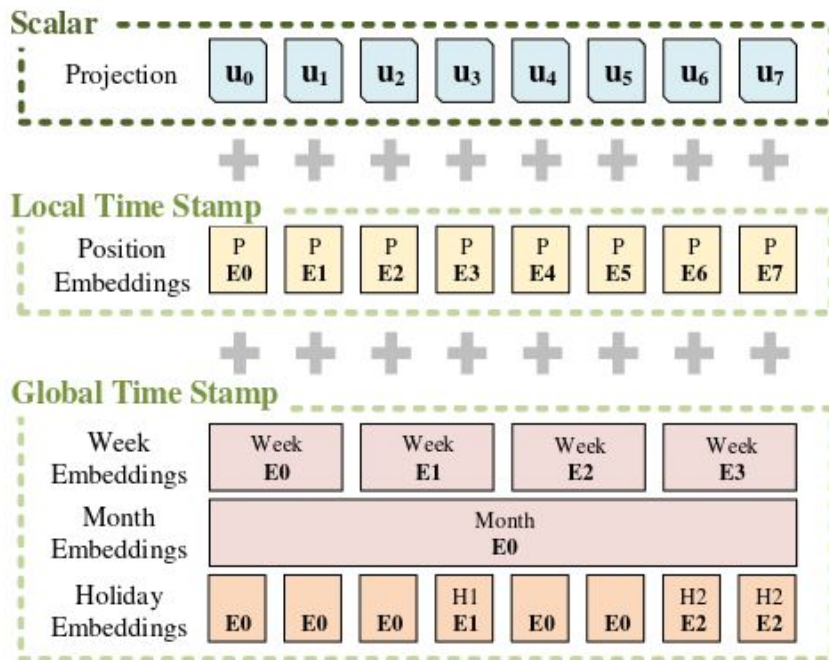


- How is stock dataset different from other dataset?
 - Stock prices are fluctuated in every day.
 - Stocks in the same sector share a similar trend even though their prices are perturbed randomly in a short-term manner
 - Especially, stocks are highly correlated in a **bull market (market index)**.
- How can we efficiently correlate multiple stocks for accurate stock movement prediction?

- Expand input dimension: concatenate market index information as input
- Multi-Level Context Aggregation: $\mathbf{h}_u^m = \mathbf{h}_u^c + \beta \mathbf{h}^i$
 - [KDD'21] Accurate Multivariate Stock Movement Prediction via Data-Axis Transformer with Multi-Level Contexts



- Embedding method
 - Informer: **Local Position Embedding + Global Learnable Stamp Embedding**



Results


- Hyperparameter
 - input sequence length of encoder: 64
 - start token length: 40
 - prediction sequence length: 5
 - num of encoder layers: 5,
 - num of decoder layers: 1,
 - training epoch: 20
 - loss function: MSE
 - learning rate: $1e-4$ (decaying to $0.7x$ every epoch)
 - batch size: 32
 - early stopping patience: 5 (If validation data loss isn't smaller than before five times, then stop.)
 - dropout: 0.05

- Predicted Pfizer's (PFE) stock using Johnson & Johnson's (JNJ), with a negative beta of -0.1128, indicating an inverse relationship due to industry competition.

Predict PFE with NASDAQ




Predict PFE with JNJ




- Comparative analysis between the stocks and the overall market index (NASDAQ).

1. JPMorgan (JPM) for Finance

Predict JPM




Predict JPM with NASDAQ




- Comparative analysis between the stocks and the overall market index (NASDAQ).

2. Drive Shack (DS) for Entertainment

Predict DS



Predict DS with NASDAQ



- Comparative analysis between the stocks and the overall market index (NASDAQ).

3. ExxonMobil (XOM) for Energy

Predict XOM




Predict XOM with NASDAQ



- Comparative analysis between the stocks and the overall market index (NASDAQ).

4. Pfizer (PFE) for Pharmaceuticals

Predict PFE



Predict PFE with NASDAQ

